

Introduction:

Hash Tables

Data Structures and Algorithms

Outline

- 1 Applications of Hashing
- 2 IP Addresses
- 3 Direct Addressing
- 4 List-based Mapping
- 5 Hash Functions
- 6 Chaining
- 7 Hash Tables

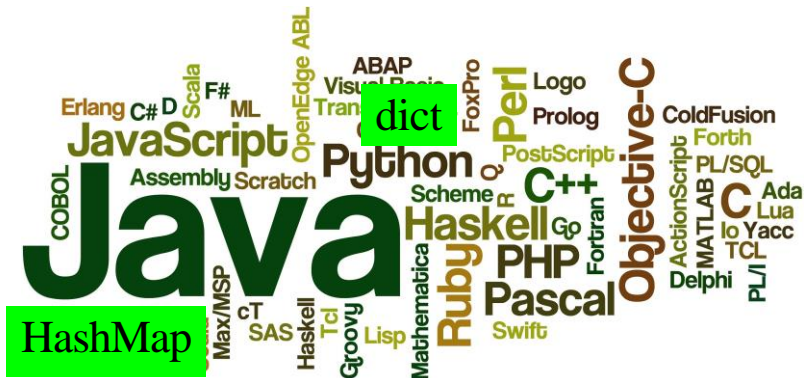
Programming Languages



Programming Languages



Programming Languages

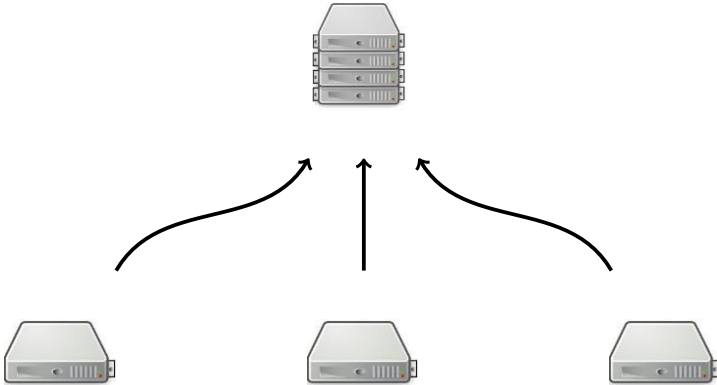


for, if, while, int

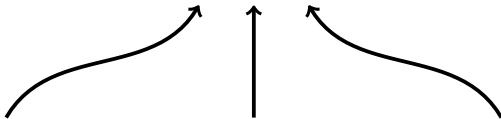
Outline

- 1 Applications of Hashing
- 2 IP Addresses
- 3 Direct Addressing
- 4 List-based Mapping
- 5 Hash Functions
- 6 Chaining
- 7 Hash Tables

Web Service



Web Service



173.194.71.102

69.171.230.68

91.210.105.134

Web Service

$$2^{32} = 4294967296$$

IP addresses



173.194.71.102

69.171.230.68

91.210.105.134

Web Service

$2^{32} = 4294967296$
IP addresses

2^{128} IPv6 addresses
number with 39 digits!



173.194.71.102

69.171.230.68

91.210.105.134

Access Log

Date	Time	IP address
09 Dec 2015	00:45:13	173.194.71.102
09 Dec 2015	00:45:15	69.171.230.68
...
...
09 Dec 2015	01:45:13	91.210.105.134

IP Access List

Analyse the access log and quickly answer queries: did anybody access the service from this *IP* during the last hour? How many times? How many *IPs* were used to access the service during the last hour?

Log Processing

- 1h of logs can contain millions of lines

Log Processing

- 1h of logs can contain millions of lines
- Too slow to process that for each query

Log Processing

- 1h of logs can contain millions of lines
- Too slow to process that for each query
- Keep count: how many times each IP appears in the last 1h of the access log

Log Processing

- 1h of logs can contain millions of lines
- Too slow to process that for each query
- Keep count: how many times each IP appears in the last 1h of the access log
- C is some data structure to store the mapping from IPs to counters

Log Processing

- 1h of logs can contain millions of lines
- Too slow to process that for each query
- Keep count: how many times each IP appears in the last 1h of the access log
- C is some data structure to store the mapping from IPs to counters
- We will learn later how to implement C

Log Processing

Time	IP address
00:45:13	173.194.71.102
00:45:13	69.171.230.68
...	...
01:45:13	173.194.71.102
01:45:13	91.210.105.134

Log Processing

	Time	IP address
	00:45:13	173.194.71.102
	00:45:13	69.171.230.68

Now	01:45:13	173.194.71.102
	01:45:13	91.210.105.134

Log Processing

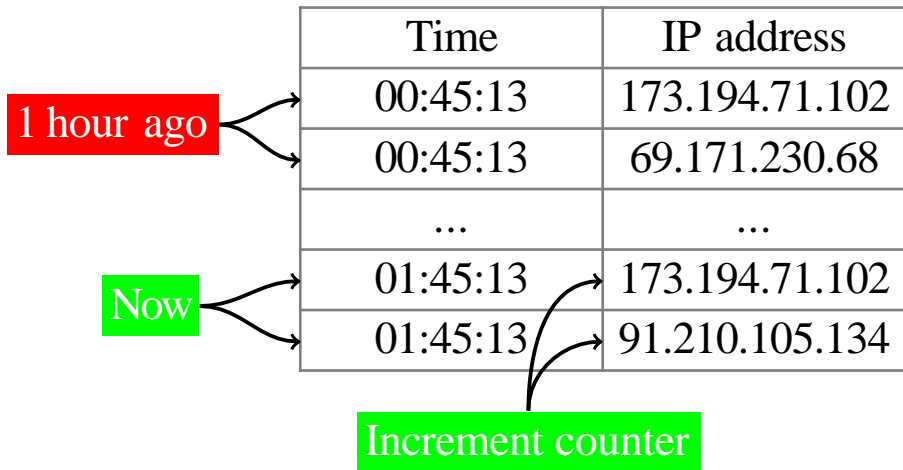
Time	IP address
00:45:13	173.194.71.102
00:45:13	69.171.230.68
...	...
01:45:13	173.194.71.102
01:45:13	91.210.105.134

Now

Increment counter

The diagram illustrates a log processing step. A table contains log entries with 'Time' and 'IP address' columns. A green box labeled 'Now' points to the 'Time' column of the last two rows, which both show '01:45:13'. A green box labeled 'Increment counter' points to the 'IP address' column of the last two rows, which show '173.194.71.102' and '91.210.105.134' respectively. This suggests that the counter is being incremented for each new IP address encountered at the current time.

Log Processing



Log Processing

Decrement counter

1 hour ago

Time	IP address
00:45:13	173.194.71.102
00:45:13	69.171.230.68
...	...
01:45:13	173.194.71.102
01:45:13	91.210.105.134

Now

Increment counter

Coming Next

How to implement the mapping \mathcal{C} ?

Outline

- 1 Applications of Hashing
- 2 IP Addresses
- 3 Direct Addressing
- 4 List-based Mapping
- 5 Hash Functions
- 6 Chaining
- 7 Hash Tables

Direct Addressing

- Need a data structure for C

Direct Addressing

- Need a data structure for C
- There are 2^{32} different IP(v4) addresses

Direct Addressing

- Need a data structure for C
- There are 2^{32} different IP(v4) addresses
- Convert IP to 32-bit integer

Direct Addressing

- Need a data structure for C
- There are 2^{32} different IP(v4) addresses
- Convert IP to 32-bit integer
- Create an integer array A of size 2^{32}

Direct Addressing

- Need a data structure for C
- There are 2^{32} different IP(v4) addresses
- Convert IP to 32-bit integer
- Create an integer array A of size 2^{32}
- Use $A[\text{int}(IP)]$ as $C[IP]$

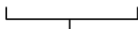
int(IP)

An IPv4 address (dotted-decimal notation)

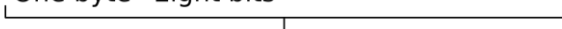
172 . 16 . 254 . 1



10101100.00010000.11111110.00000001



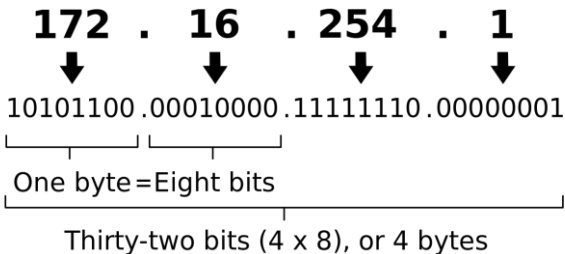
One byte=Eight bits



Thirty-two bits (4 x 8), or 4 bytes

int(IP)

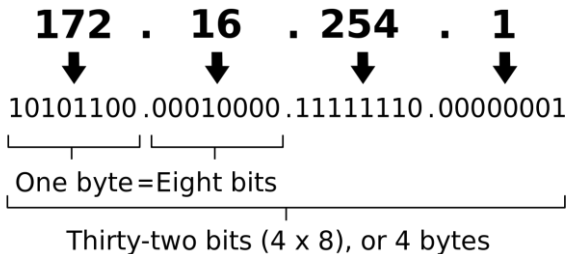
An IPv4 address (dotted-decimal notation)



■ $\text{int}(0.0.0.1) = 1$

int(IP)

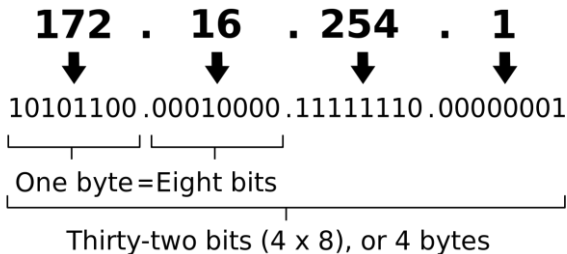
An IPv4 address (dotted-decimal notation)



- $\text{int}(0.0.0.1) = 1$
- $\text{int}(172.16.254.1) = 2886794753$

int(IP)

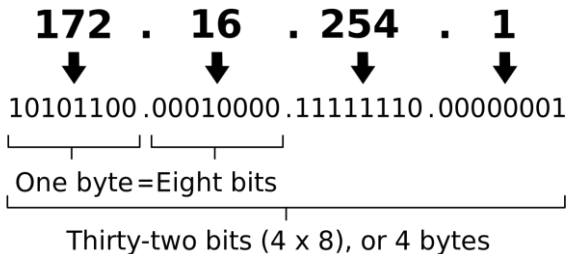
An IPv4 address (dotted-decimal notation)



- $\text{int}(0.0.0.1) = 1$
- $\text{int}(172.16.254.1) = 2886794753$
- $\text{int}(69.171.230.68) =$

int(IP)

An IPv4 address (dotted-decimal notation)



- $\text{int}(0.0.0.1) = 1$
- $\text{int}(172.16.254.1) = 2886794753$
- $\text{int}(69.171.230.68) = 1168893508$

Asymptotics

- UpdateAccessList is $O(1)$ per log line

Asymptotics

- UpdateAccessList is $O(1)$ per log line
- AccessedLastHour is $O(1)$

Asymptotics

- UpdateAccessList is $O(1)$ per log line
- AccessedLastHour is $O(1)$
- But need 2^{32} memory even for few IPs

Asymptotics

- UpdateAccessList is $O(1)$ per log line
- AccessedLastHour is $O(1)$
- But need 2^{32} memory even for few IPs
- IPv6: 2^{128} won't fit in memory

Asymptotics

- UpdateAccessList is $O(1)$ per log line
- AccessedLastHour is $O(1)$
- But need 2^{32} memory even for few IPs
- IPv6: 2^{128} won't fit in memory
- In general: $O(N)$ memory, $N = |S|$

Outline

- 1 Applications of Hashing
- 2 IP Addresses
- 3 Direct Addressing
- 4 Hash Functions
- 5 Chaining
- 6 Hash Tables

Encoding IPs

- Encode IPs with small numbers

Encoding IPs

- Encode IPs with small numbers
- I.e. numbers from 0 to 999

Encoding IPs

- Encode IPs with small numbers
- i.e. numbers from 0 to 999
- Different codes for currently active IPs

Hash Function

Definition

For any set of objects S and any integer $m > 0$, a function $h : S \rightarrow \{0, 1, \dots, m - 1\}$ is called a **hash function**.

Hash Function

Definition

For any set of objects S and any integer $m > 0$, a function $h : S \rightarrow \{0, 1, \dots, m - 1\}$ is called a **hash function**.

Definition

m is called the **cardinality** of hash function h .

Desirable Properties

- h should be fast to compute

Desirable Properties

- h should be fast to compute
- Different values for different objects

Desirable Properties

- h should be fast to compute
- Different values for different objects
- Direct addressing with $O(m)$ memory

Desirable Properties

- h should be fast to compute
- Different values for different objects
- Direct addressing with $O(m)$ memory
- Want small cardinality m

Desirable Properties

- h should be fast to compute
- Different values for different objects
- Direct addressing with $O(m)$ memory
- Want small cardinality m
- Impossible to have all different values if number of objects $|S|$ is more than m

Popular Hash Function

Division method: Choose a number m larger than the number n of keys in K . (The number m is usually chosen to be a prime number or a number without small divisors, since this frequently minimizes the number of collisions.) The hash function H is defined by

$$H(k) = k \text{ (mod } m) \text{ or } H(k) = k \text{ (mod } m) + 1$$

Here $k \text{ (mod } m)$ denotes the remainder when k is divided by m . The second formula is used when we want the hash addresses to range from 1 to m rather than from 0 to $m - 1$.

Popular Hash Function

Midsquare method: The key k is squared. Then the hash function H is defined by

$$H(k) = l$$

where l is obtained by deleting digits from both ends of k^2 . We emphasize that the same positions of k^2 must be used for all of the keys.

Popular Hash Function

Folding method: The key k is partitioned into a number of parts, k_1, \dots, k_r , where each part, except possibly the last, has the same number of digits as the required address. Then the parts are added together, ignoring the last carry. That is,

$$H(k) = k_1 + k_2 + \dots + k_r$$

where the leading-digit carries, if any, are ignored. Sometimes, for extra “milling,” the even-numbered parts, k_2, k_4, \dots , are each reversed before the addition.

Question

Suppose a company with 68 employees assigns a 4-digit employee number to each employee which is used as the primary key in the company's employee file. Suppose L consists of 100 two-digit addresses: 00, 01, 02, ..., 99. Compute the locations to which the keys 3205, 7148, and 2345 are mapped.

(a) Division method. Choose a prime number m close to 99, such as $m = 97$.

Then

$$H(3205) = 4, H(7148) = 67, H(2345) = 17$$

That is, dividing 3205 by 97 gives a remainder of 4, dividing 7148 by 97 gives a remainder of 67, and dividing 2345 by 97 gives a remainder of 17. In the case that the memory addresses begin with 01 rather than 00, we choose that the function $H(k) = k(\bmod m) + 1$ to obtain:

$$H(3205) = 4 + 1 = 5, H(7148) = 67 + 1 = 68, H(2345) = 17 + 1 = 18$$

(b) Mid-square method. The following calculations are performed:

k :	3205	7148	2345
k^2 :	10 272 025	51 093 904	5 499 025
$H(k)$:	72	93	99

Observe that the fourth and fifth digits, counting from the right, are chosen for the hash address.

Question

Suppose a company with 68 employees assigns a 4-digit employee number to each employee which is used as the primary key in the company's employee file. Suppose L consists of 100 two-digit addresses: 00, 01, 02, ..., 99. Compute the locations to which the keys 3205, 7148, and 2345 are mapped.

(c) *Folding method.* Chopping the key k into two parts and adding yields the following hash addresses:

$$H(3205) = 32 + 05 = 37, H(7148) = 71 + 48 = 19, H(2345) = 23 + 45 = 68$$

Observe that the leading digit 1 in $H(7148)$ is ignored. Alternatively, one may want to reverse the second part before adding, thus producing the following hash addresses:

$$H(3205) = 32 + 50 = 82, H(7148) = 71 + 84 + 55, H(2345) = 23 + 54 = 77$$

Question

Consider a hash table of size $m=1000$ and a corresponding hash function $h(k)=\lfloor m(kA \bmod 1) \rfloor$ for $A=(\sqrt{5}-1)/2$. Compute the locations to which the keys 6161, 6262, 6363, 6464, and 6565 are mapped.

- $h(61) = \lfloor 1000(61 \cdot \frac{\sqrt{5}-1}{2} \bmod 1) \rfloor = 700.$
- $h(62) = \lfloor 1000(62 \cdot \frac{\sqrt{5}-1}{2} \bmod 1) \rfloor = 318.$
- $h(63) = \lfloor 1000(63 \cdot \frac{\sqrt{5}-1}{2} \bmod 1) \rfloor = 936.$
- $h(64) = \lfloor 1000(64 \cdot \frac{\sqrt{5}-1}{2} \bmod 1) \rfloor = 554.$
- $h(65) = \lfloor 1000(65 \cdot \frac{\sqrt{5}-1}{2} \bmod 1) \rfloor = 172.$

Collisions

Definition

When $h(o_1) = h(o_2)$ and $o_1 \neq o_2$, this is a collision.

Outline

- 1 Applications of Hashing
- 2 IP Addresses
- 3 Direct Addressing
- 4 List-based Mapping
- 5 Hash Functions
- 6 Chaining
- 7 Hash Tables

Map

Store mapping from objects to other objects:

- Filename → location of the file on disk
- Student ID → student name
- Contact name → contact phone number

Map

Store mapping from objects to other objects:

- Filename \rightarrow location of the file on disk
- Student ID \rightarrow student name
- Contact name \rightarrow contact phone number

Definition

Map from S to V is a data structure with methods $\text{HasKey}(O)$, $\text{Get}(O)$, $\text{Set}(O, v)$, where $O \in S$, $v \in V$.

Chaining

0
1
2
3
4
5
6
7

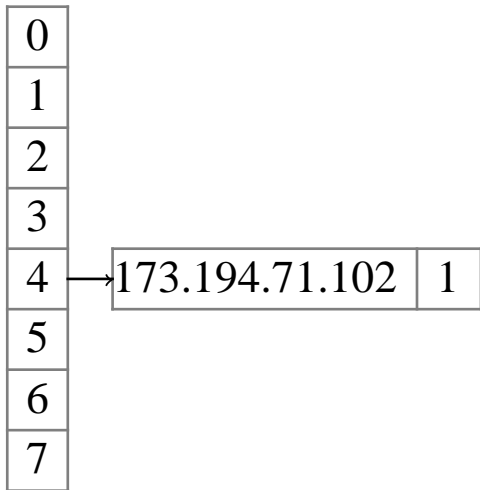
Chaining

$$h(173.194.71.102) = 4$$

0
1
2
3
4
5
6
7

Chaining

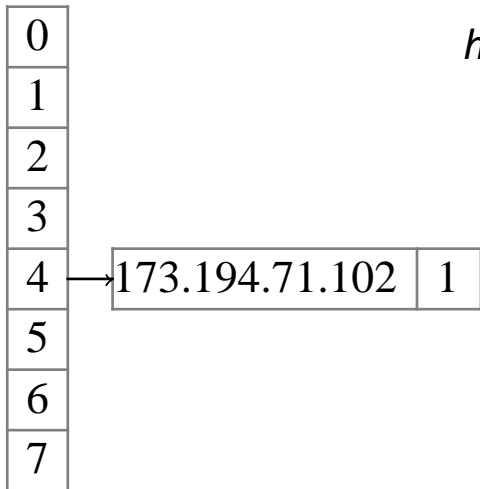
$$h(173.194.71.102) = 4$$



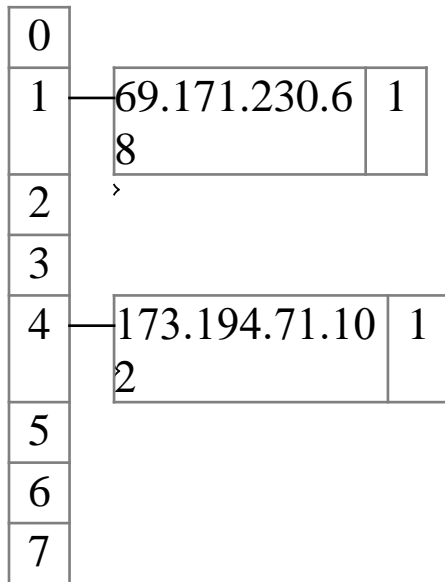
Chaining

$$h(173.194.71.102) = 4$$

$$h(69.171.230.68) = 1$$



Chaining



$$h(173.194.71.102) = 4$$

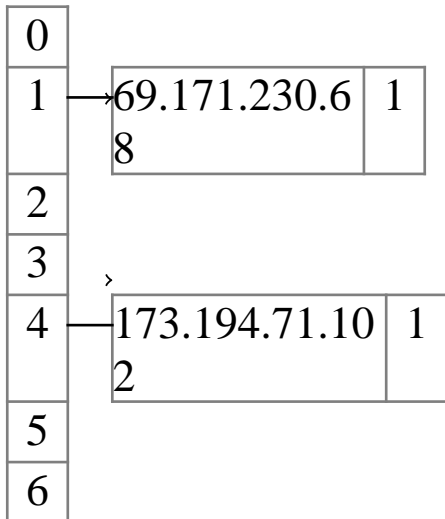
$$h(69.171.230.68) = 1$$

Chaining

$$h(173.194.71.102) = 4$$

$$h(69.171.230.68) = 1$$

$$h(173.194.71.102) = 4$$

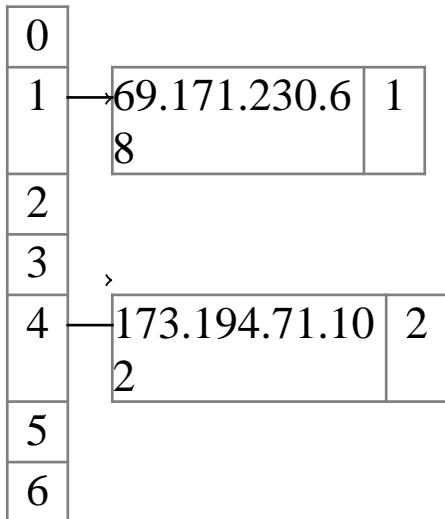


Chaining

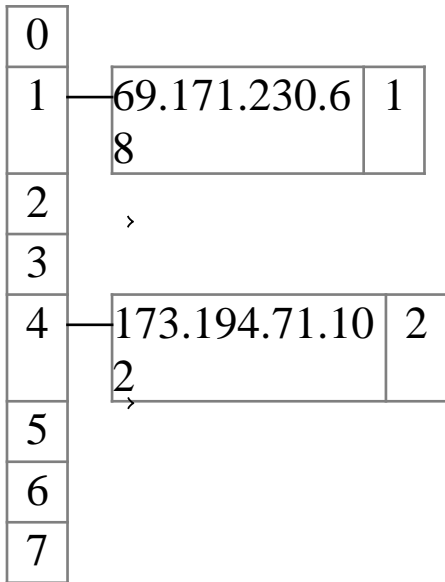
$$h(173.194.71.102) = 4$$

$$h(69.171.230.68) = 1$$

$$h(173.194.71.102) = 4$$



Chaining



$$h(173.194.71.102) = 4$$

$$h(69.171.230.68) = 1$$

$$h(173.194.71.102) = 4$$

$$h(91.210.105.134) = 4$$

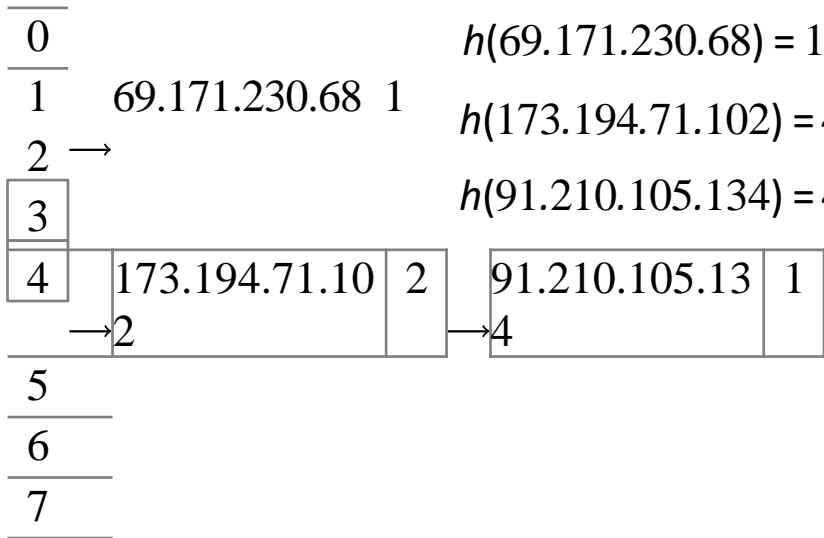
Chaining

$$h(173.194.71.102) = 4$$

$$h(69.171.230.68) = 1$$

$$h(173.194.71.102) = 4$$

$$h(91.210.105.134) = 4$$



Chains — array of chains

Each chain is a list of pairs (object, value)

HasKey(object)

```
chain ← Chains[hash(object)]  
for (key, value) in chain:  
    if key == object:  
        return true  
return false
```

Get(object)

```
chain ← Chains[hash(object)]  
for (key, value) in chain:  
    if key == object:  
        return value  
return N/A
```


Set(object, value)

```
chain ← Chains[hash(object)]  
for pair in chain:  
    if pair.key == object:  
        pair.value ← value  
    return  
chain.Append((object, value))
```

Lemma

Let c be the length of the longest chain in *chains*. Then the running time of HasKey, Get, Set is $\Theta(c + 1)$.

Lemma

Let c be the length of the longest chain in A . Then the running time of HasKey, Get, Set is $\Theta(c + 1)$.

Proof

- If $L = A[h(O)]$, $\text{len}(L) = c$, $O \notin L$, need to scan all c items

Lemma

Let c be the length of the longest chain in A . Then the running time of HasKey, Get, Set is $\Theta(c + 1)$.

Proof

- If $L = A[h(O)]$, $\text{len}(L) = c$, $O \notin L$, need to scan all c items
- If $c = 0$, we still need $O(1)$ time



Lemma

Let n be the number of different keys O currently in the map and m be the cardinality of the hash function. Then the memory consumption for chaining is $\Theta(n + m)$.

Lemma

Let n be the number of different keys O currently in the map and m be the cardinality of the hash function. Then the memory consumption for chaining is $\Theta(n + m)$.

Proof

- $\Theta(n)$ to store n pairs (O, v)

Lemma

Let n be the number of different keys O currently in the map and m be the cardinality of the hash function. Then the memory consumption for chaining is $\Theta(n + m)$.

Proof

- $\Theta(n)$ to store n pairs (O, v)
- $\Theta(m)$ to store array A of size m



Outline

- 1 Applications of Hashing
- 2 IP Addresses
- 3 Direct Addressing
- 4 List-based Mapping
- 5 Hash Functions
- 6 Chaining
- 7 Hash Tables

Set

Definition

Set is a data structure with methods
 $\text{Add}(O)$, $\text{Remove}(O)$, $\text{Find}(O)$.

Set

Definition

Set is a data structure with methods
 $\text{Add}(O)$, $\text{Remove}(O)$, $\text{Find}(O)$.

Examples

- IPs accessed during last hour

Set

Definition

Set is a data structure with methods
 $\text{Add}(O)$, $\text{Remove}(O)$, $\text{Find}(O)$.

Examples

- IPs accessed during last hour
- Students on campus

Set

Definition

Set is a data structure with methods
 $\text{Add}(O)$, $\text{Remove}(O)$, $\text{Find}(O)$.

Examples

- IPs accessed during last hour
- Students on campus
- Keywords in a programming language

Implementing Set

Two ways to implement a set using chaining:

- Set is equivalent to map from S to

$$V = \{true, false\}$$

Implementing Set

Two ways to implement a set using chaining:

- Set is equivalent to map from S to $V = \{true, false\}$
- Store just objects O instead of pairs (O, v) in chains

$h : S \rightarrow \{0, 1, \dots, m - 1\}$

$O, O' \in S$

$A \leftarrow$ array of m lists (chains) of objects O

Find(O)

$L \leftarrow A[h(O)]$

for O' in L :

 if $O' == O$:

 return true

return false

Add(O)

```
 $L \leftarrow A[h(O)]$   
for  $O'$  in  $L$ :  
    if  $O' == O$ :  
        return  
 $L.Append(O)$ 
```


Remove(O)

if not Find(O):

 return

$L \leftarrow A[h(O)]$

$L.\text{Erase}(O)$

Hash Table

Definition

An implementation of a set or a map using hashing is called a hash table.

Programming Languages

Set:

- `unordered_set` in C++
- `HashSet` in Java
- `set` in Python

Map:

- `unordered_map` in C++
- `HashMap` in Java
- `dict` in Python

Conclusion

- Chaining is a technique to implement a hash table

Conclusion

- Chaining is a technique to implement a hash table
- Memory consumption is $O(n + m)$

Conclusion

- Chaining is a technique to implement a hash table
- Memory consumption is $O(n + m)$
- Operations work in time $O(c + 1)$

Conclusion

- Chaining is a technique to implement a hash table
- Memory consumption is $O(n + m)$
- Operations work in time $O(c + 1)$
- How to make both m and c small?