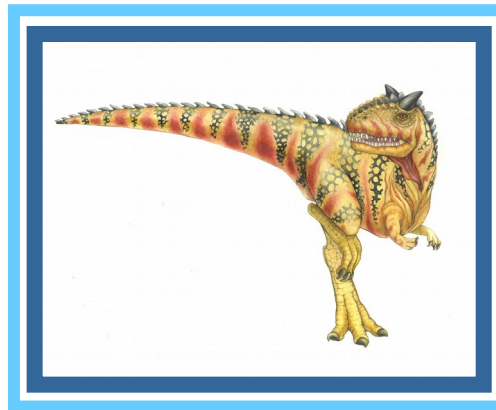# Chapter 2: Operating-System Structures

Objectives

- To describe the services an operating system provides to users, processes, and other systems

- To discuss the various ways of structuring an operating system

- To explain how operating systems are installed and customized and how they boot

Operating System Services

- Operating systems provide an environment for the execution of programs.

- Operating systems provides certain services to:
  - Programs
  - Users of those programs

- Basically two types of services:
  - Set of operating-system services provides functions that are helpful to the user:
  - Set of operating-system functions for ensuring the efficient operation of the system itself via resource sharing

OS Services Helpful to the User

- **User interface** - Almost all operating systems have a **user interface** (**UI**). This interface can take several forms:

  - **Command-Line** (**CLI**) -- uses text commands and a method for entering them (say, a keyboard for typing in commands in a specific format with specific options).

  - **Graphics User Interface** (**GUI**) -- the interface is a window system with a pointing device to direct I\O, choose from menus, and make selections and a keyboard to enter text..

  - **Batch Interface** -- commands and directives to control those commands are entered into files, and those files are executed

  Some systems provide two or all three of these variations.

- **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)

- **I/O operations** - A running program may require I/O, which may involve a file or an I/O device

OS Services Helpful to the User (Cont.)

- **File-system manipulation** - Programs need to read and write files and directories, create and delete them, search them, list file Information, permission management.

- **Communications** – Processes may exchange information, on the same computer or between computers over a network
  - Communications may be via shared memory or through message passing (packets moved by the OS)

- **Error detection** – OS needs to be constantly aware of possible errors
  - May occur in the CPU and memory hardware, in I/O devices, in user program
  - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
  - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

OS Services for Ensuring Efficient Operation

- **Resource allocation -** When multiple users or multiple jobs are running concurrently, resources must be allocated to each of them

  - Many types of resources - CPU cycles, main memory, file storage, I/O devices.

- **Accounting -** To keep track of which users use how much and what kinds of computer resources

- **Protection and security -** The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other

  - **Protection** involves ensuring that all access to system resources is controlled

  - **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts
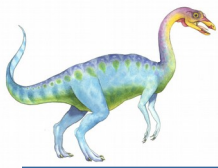
A View of Operating System Services
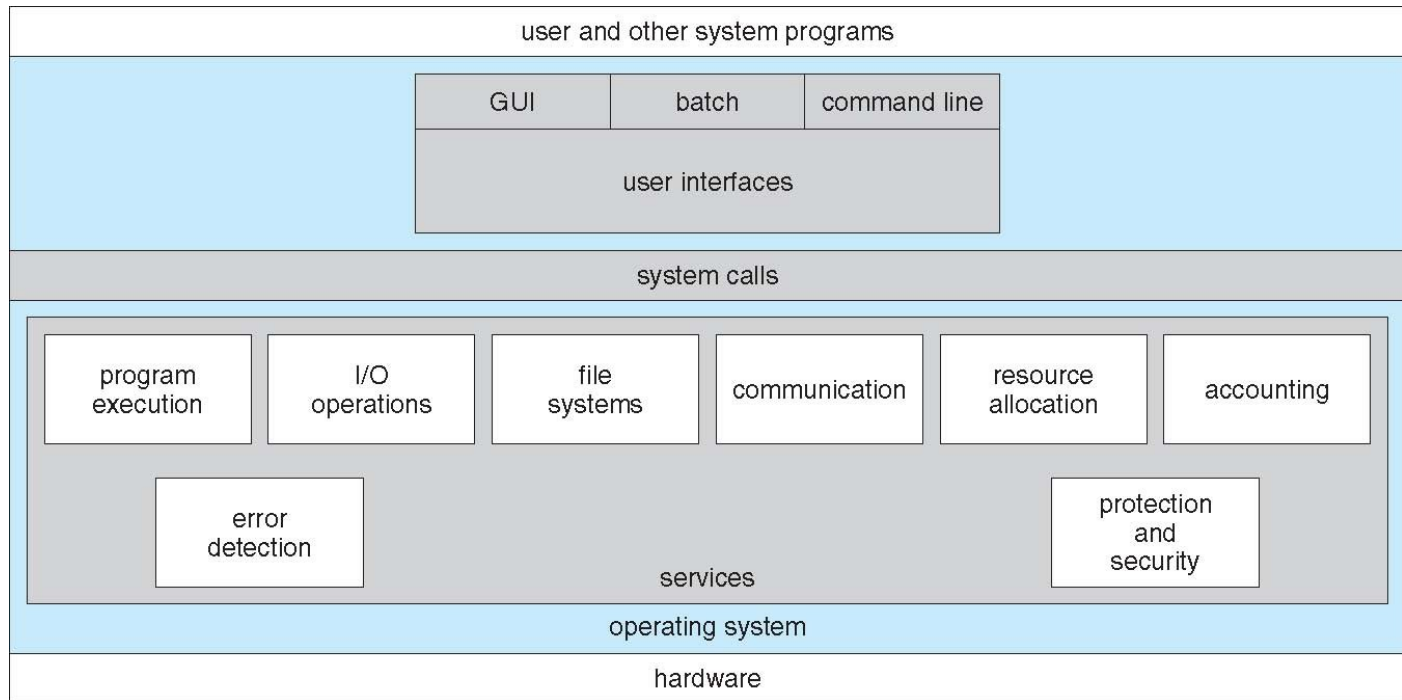
Command Interpreters (CLI)

CLI allows users to directly enter commands to be performed by the operating system.

- Some operating systems include the command interpreter in the kernel.

- Some operating systems, such as Windows and UNIX, treat the command interpreter as a special program that is running when a job is initiated or when a user first logs on.

- On systems with multiple command interpreters to choose from, the interpreters are known as **shells.**

- The main function of the command interpreter is to get and execute the next user-specified command.

- Sometimes commands built-in, sometimes just names of programs
 - If the latter, adding new features doesn't require shell modification

The Bourne shell command interpreter in Solaris

Bourne Shell Command Interpreter

Graphical User Interfaces (GUI)

- User-friendly **desktop** metaphor interface
  - Usually mouse, keyboard, and monitor
  - **Icons** represent files, programs, actions, etc
  - Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a **folder**)
  - Invented at Xerox PARC
- Many systems now include both CLI and GUI interfaces
  - Microsoft Windows is GUI with CLI "command" shell
  - Apple Mac OS X is "Aqua" GUI interface with UNIX kernel underneath and shells available
  - Unix and Linux have CLI with optional GUI interfaces (CDE, KDE, GNOME)
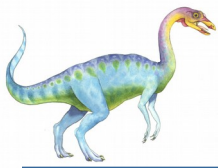
Touchscreen Interfaces

- Touchscreen devices require new interfaces
  - Mouse not possible or not desired
  - Actions and selection based on gestures
  - Virtual keyboard for text entry
- Voice commands.

The Mac OS X GUI

System Calls

- Programming interface to the services provided by the OS

- Typically written in a high-level language (C or C++)

- Mostly accessed by programs via a high-level **Application Programming Interface** (**API**) rather than direct system call

- Three most common APIs are:

  - Win32 API for Windows,

  - POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X),

  - Java API for the Java virtual machine (JVM)

Note that the system-call names used throughout this text are generic

Example of System Calls

■ System call sequence to copy the contents of one file to another file

source file → destination file

Example System Call Sequence

Acquire input file name
  Write prompt to screen
  Accept input
Acquire output file name
  Write prompt to screen
  Accept input
Open the input file
  if file doesn't exist, abort
Create output file
  if file exists, abort
Loop
  Read from input file
  Write to output file
Until read fails
Close output file
Write completion message to screen
Terminate normally

Example of Standard API

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the `man` page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

```
#include <unistd.h>

ssize_t       read(int fd, void *buf, size_t count)
```

| return value | function name | parameters |

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns −1.

System Call Implementation

- Typically, a number is associated with each system call
  - **System-call interface** maintains a table indexed according to these numbers
- The system call interface invokes the intended system call in OS kernel and returns status of the system call and any return values
- The caller need not know a thing about how the system call is implemented
  - Just needs to obey the API and understand what the OS will do as a result call
  - Most details of OS interface hidden from programmer by API
    - Managed by run-time support library (set of functions built into libraries included with compiler)

System Call -- OS Relationship

The handling of a user application invoking the open() system call
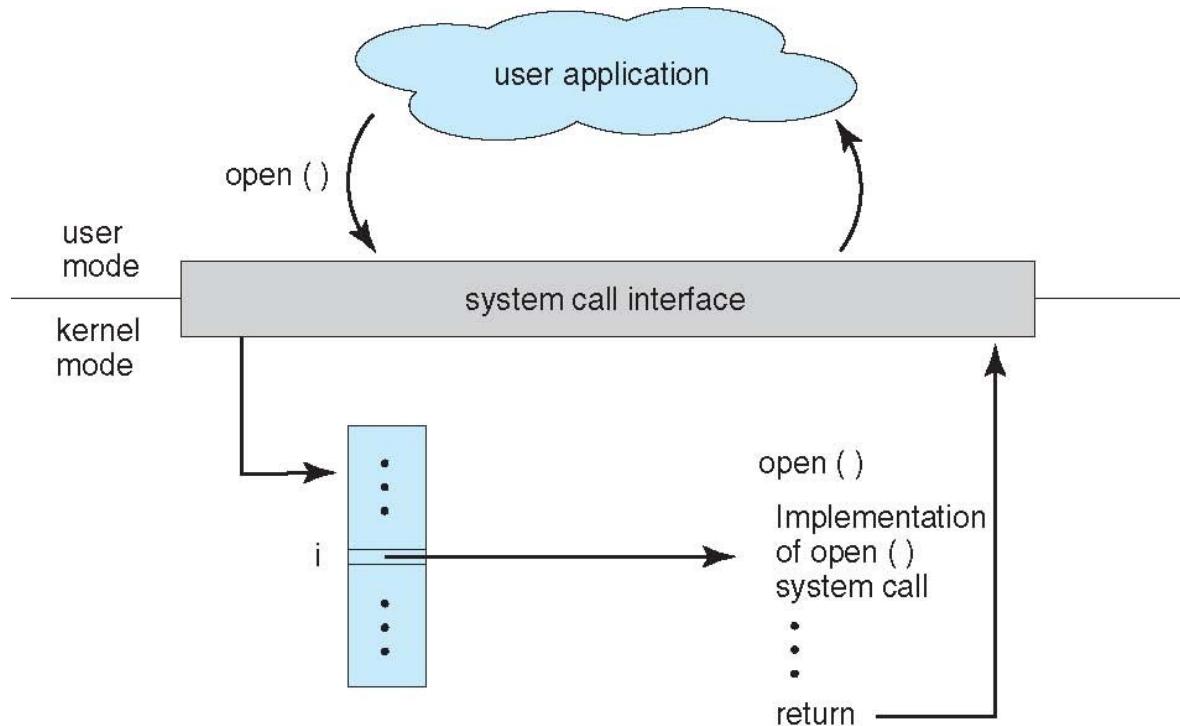
System Call Parameter Passing

- Often, more information is required than simply identity of desired system call
  - Exact type and amount of information vary according to OS and call

- Three general methods used to pass parameters to the OS
  - Simplest: pass the parameters in registers
    - In some cases, may be more parameters than registers
  - Parameters stored in a block, or table, in memory, and address of block passed as a parameter in a register
    - This approach taken by Linux and Solaris
  - Parameters placed, or **pushed**, onto the **stack** by the program and **popped** off the stack by the operating system
  - Block and stack methods do not limit the number or length of parameters being passed

Parameter Passing via Table

x points to a block of parameters. x is loaded into a register
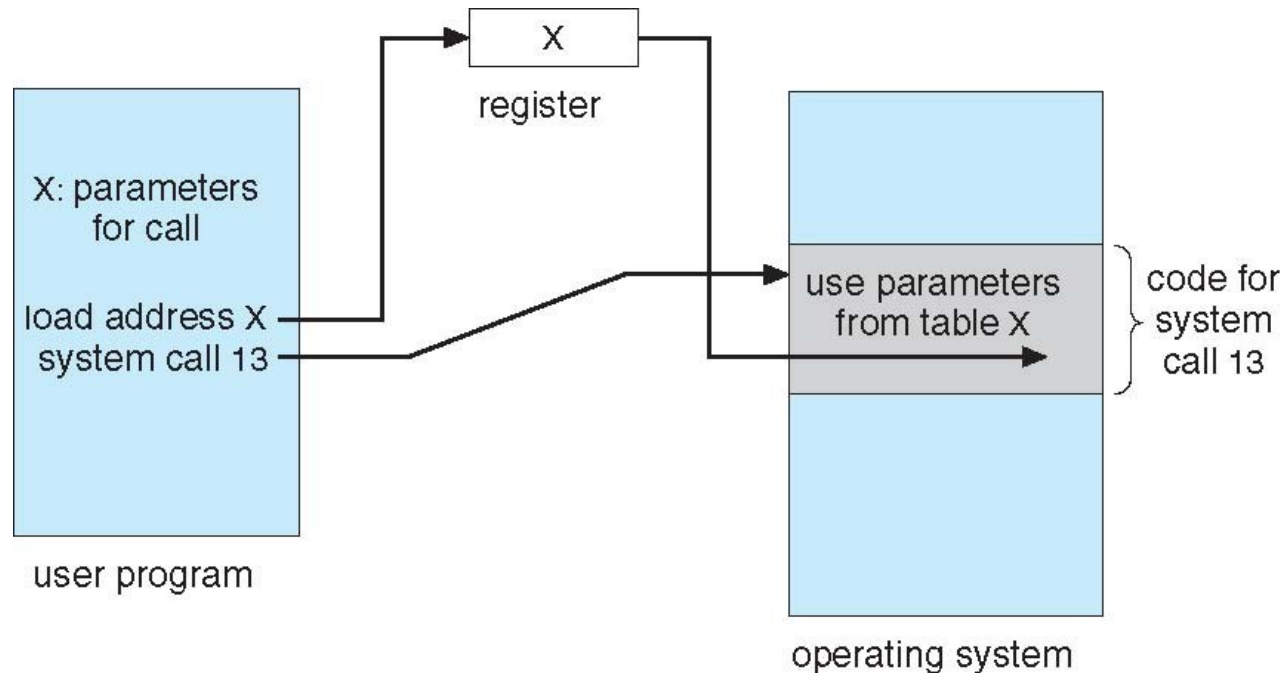
Types of System Calls

- System calls can be grouped roughly into six major categories:
    - Process control,
    - File manipulation,
    - Device manipulation,
    - Information maintenance,
    - Communications,
    - Protection.
- The figure in the slide # 28 summarizes the types of system calls normally provided by an operating system.

System Calls – Process Control

- create process, terminate process

- end, abort

- load, execute

- get process attributes, set process attributes

- wait for time

- wait event, signal event

- allocate and free memory

- Dump memory if error

- **Debugger** for determining **bugs, single step** execution

- **Locks** for managing access to shared data between processes

System Calls – File Management

- Create file
- Delete file
- Open and Close file
- Read, Write, Reposition
- Get and Set file attributes

System Calls – Device Management

- request device, release device

- read, write, reposition

- get device attributes, set device attributes

- logically attach or detach devices

System Calls -- Information Maintenance

- get time or date,

- set time or date

- get system data,

- set system data

- get and set process, file, or device attributes

System Calls – Communications

- create, delete communication connection

- if **message passing model:**
  - send, receive message
    - To **host name** or **process name**
    - From **client** to **server**

- If **shared-memory model:**
  - create and gain access to memory regions

- transfer status information

- attach and detach remote devices

System Calls -- Protection

- Control access to resources
- Get and set permissions
- Allow and deny user access

Examples of Windows and Unix System Calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitlializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

Example -- Standard C Library

C program invoking printf() library call, which calls write() system call
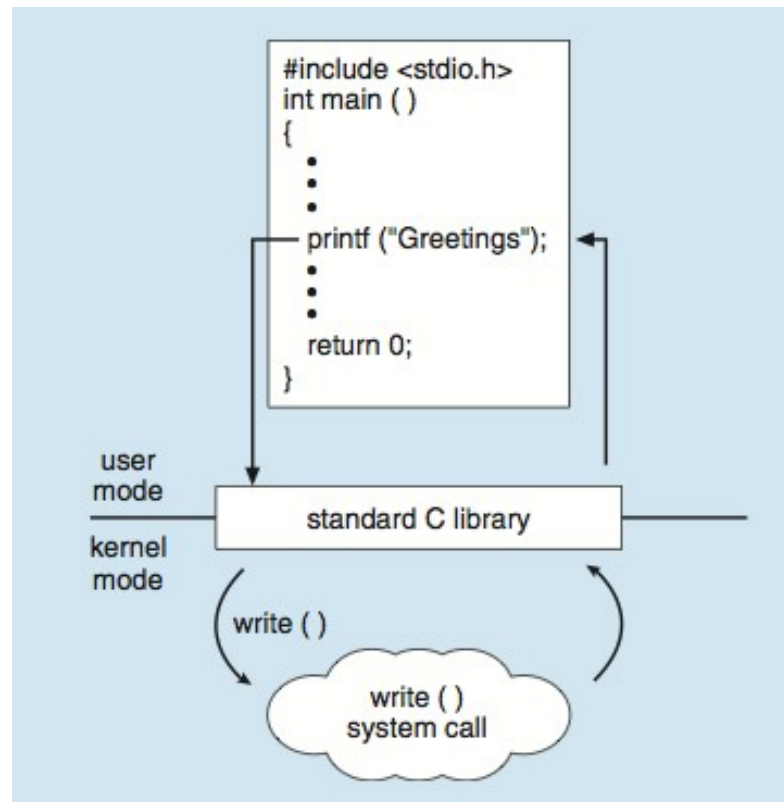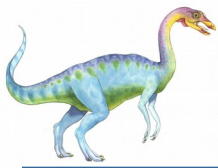
System Programs

- System programs provide a convenient environment for program development and execution.

- Some of them are simply user interfaces to system calls. Others are considerably more complex.

- They can be divided into:
  - File manipulation
  - Status information sometimes stored in a File modification
  - Programming language support
  - Program loading and execution
  - Communications
  - Background services
  - Application programs

- Most users' view of the operation system is defined by system programs, not the actual system calls

System Programs

- **File management**
  - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories

- **Status information**
  - Some programs ask the system for information - date, time, amount of available memory, disk space, number of users
  - Others programs provide detailed performance, logging, and debugging information
  - Typically, these programs format and print the output to the terminal or other output devices
  - Some systems implement a **registry** - used to store and retrieve configuration information

System Programs (Cont.)

- **File modification**
  - Text editors to create and modify files
  - Special commands to search contents of files or perform transformations of the text
- **Programming-language support** - Compilers, assemblers, debuggers and interpreters sometimes provided
- **Program loading and execution**- Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language
- **Communications** - Provide the mechanism for creating virtual connections among processes, users, and computer systems
  - Allow users to send messages to one another's screens, browse web pages, send electronic-mail messages, log in remotely, transfer files from one machine to another

System Programs (Cont.)

- **Background Services**
  - Launch at boot time
    - Some for system startup, then terminate
    - Some from system boot to shutdown
  - Provide facilities like disk checking, process scheduling, error logging, printing
  - Run in user context not kernel context
  - Known as **services**, **subsystems**, **daemons**

- **Application programs**
  - Don't pertain to system
  - Run by users
  - Not typically considered part of OS
  - Launched by command line, mouse click, finger poke

Operating System Design and Implementation

- Design and Implementation of OS not "solvable", but some approaches have proven successful

- Internal structure of different Operating Systems can vary widely

- Start the design by defining goals and specifications

- Affected by choice of hardware, type of system – batch, time sharing, single user, multiuser, distributed, real-time

- Two groups in terms of defining goals:

  - User goals – operating system should be convenient to use, easy to learn, reliable, safe, and fast

  - System goals – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient

- Specifying and designing an OS is highly creative task of **software engineering**

Mechanisms and Polices

- Important principle to separate

  Policy: *What* will be done?
  Mechanism: *How* to do it?

- Mechanisms determine how to do something, policies decide what will be done

- The separation of policy from mechanism is a very important principle, it allows maximum flexibility if policy decisions are to be changed later (example – timer)

Implementation

- Much variation
  - Early Operating Systems were written in assembly language
  - Then with system programming languages like Algol, PL/1
  - Now C, C++
- Actually usually a mix of languages
  - Lowest levels in assembly
  - Main body in C
  - Systems programs in C, C++, scripting languages like PERL, Python, shell scripts
- High-level language easier to **port** to other hardware
  - But slower
- **Emulation** can allow an OS to run on non-native hardware

Operating System Structure

- Various ways to structure an operating system:
  - Monolithic structure
    - Simple structure – MS-DOS
    - More complex – UNIX
    - More complex – Linux
  - Layered – An abstraction
  - Microkernel - Mach

MS-DOS

- MS-DOS – written to provide the most functionality in the least amount of space

- MS-DOS was limited by hardware funcionality.

  - Not divided into modules

  - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated
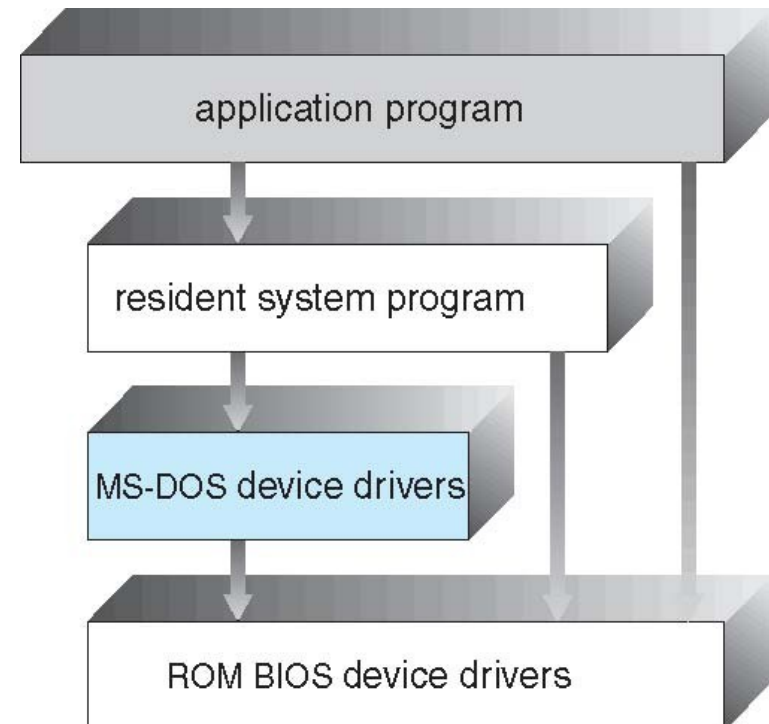
UNIX

- UNIX – the original UNIX operating system had limited structuring and was limited by hardware functionality.

- The UNIX OS consists of two separable parts

  - Systems programs

  - The kernel

    ‣ Consists of everything below the system-call interface and above the physical hardware

    ‣ Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level

Traditional UNIX System Structure

(the users)		
shells and commands compilers and interpreters system libraries		
system-call interface to the kernel		
signals terminal handling character I/O system terminal drivers	file system swapping block I/O system disk and tape drivers	CPU scheduling page replacement demand paging virtual memory
kernel interface to the hardware		
terminal controllers terminals	device controllers disks and tapes	memory controllers physical memory

Kernel

Linux System Structure

Modularity

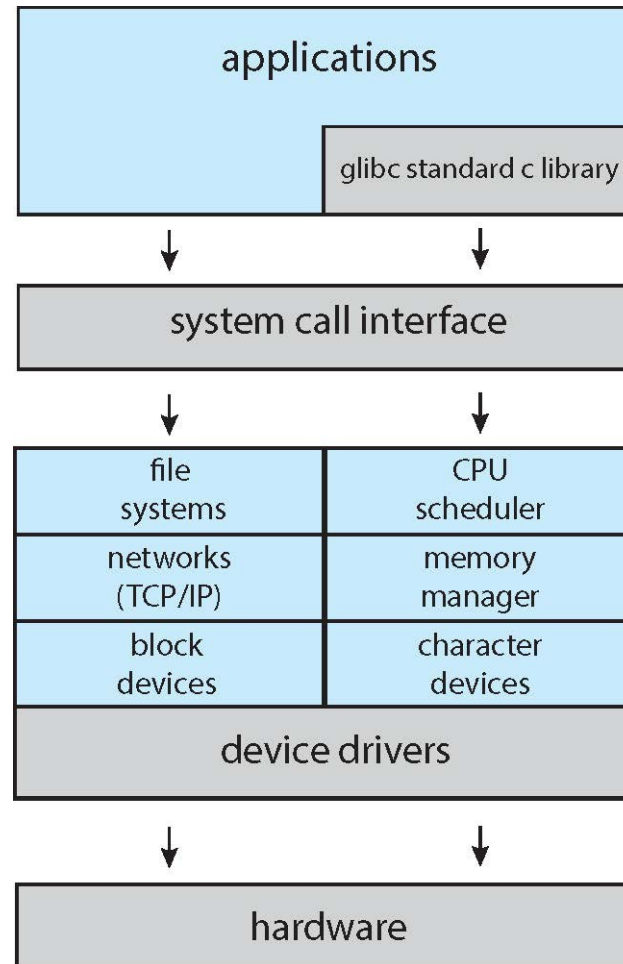- The monolithic approach results in a situation where changes to one part of the system can have wide-ranging effects to other parts.

- Alternatively, we could design system where the operating system is divided into separate, smaller components that have specific and limited functionality. The sum of all these components comprises the kernel.

- The advantage of this modular approach is that changes in one component only affect that component, and no others, allowing system implementers more freedom when changing the inner workings of the system and in creating modular operating systems.

Layered Approach

- A system can be made modular in many ways. One method is the **layered approach**, in which the operating system is broken into a number of layers (levels). The bottom layer (layer 0) is the hardware; the highest (layer N) is the user interface.

Layered Approach (Cont.)

- A typical operating-system layer -- say, layer M -- consists of data structures and a set of routines that can be invoked by higher-level layers. Layer M in turn, can invoke operations on lower-level layers.

- Each layer is implemented only with operations provided by lower-level layers. A layer does not need to know how these operations are implemented; it needs to know only what these operations do.

- The main advantage of the layered approach is simplicity of construction and debugging. The layers are selected so that each uses functions (operations) and services of only lower-level layers.

- This approach simplifies debugging and system verification. The first layer can be debugged without any concern for the rest of the system. Once the first layer is debugged, its correct functioning can be assumed while the second layer is debugged, and so on. If an error is found during the debugging of a particular layer, the error must be on that layer.

Microkernel System Structure

- Moves as much from the kernel into user space

- **Mach** example of **microkernel**

  - Mac OS X kernel (**Darwin**) partly based on Mach

- Communication takes place between user modules using **message passing**

- Benefits:

  - Easier to extend a microkernel

  - Easier to port the operating system to new architectures

  - More reliable (less code is running in kernel mode)

  - More secure

- Detriments:

  - Performance overhead of user space to kernel space communication

Microkernel System Structure
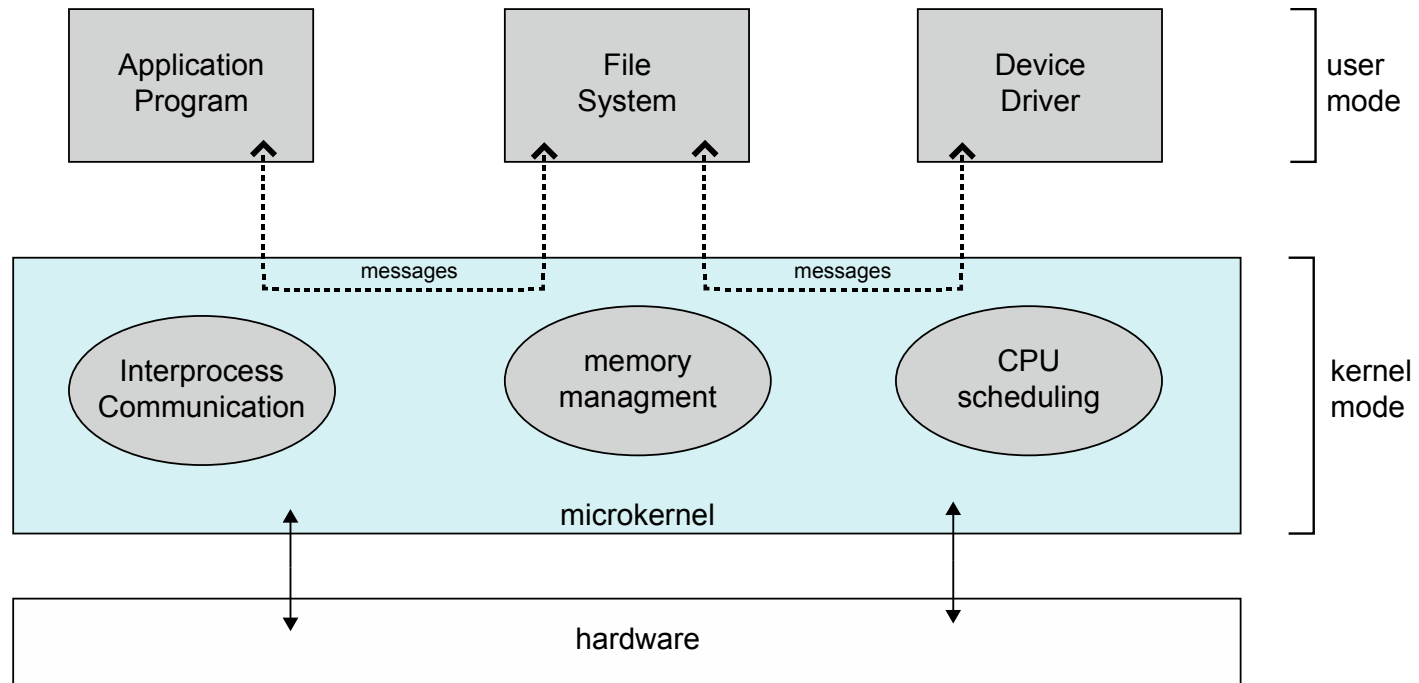
Modules

- Many modern operating systems implement **loadable kernel modules**

  - Uses object-oriented approach
  - Each core component is separate
  - Each talks to the others over known interfaces
  - Each is loadable as needed within the kernel
- Overall, similar to layers but with more flexible
  - Linux, Solaris, etc

Hybrid Systems

- In practice, very few operating systems adopt a single, strictly defined structure.

- Instead, they combine different structures, resulting in hybrid systems that address performance, security, and usability issues.

  - For example, Linux is monolithic, because having the operating system in a single address space provides very efficient performance. However, Linux are also modular, so that new functionality can be dynamically added to the kernel.

- We cover the structure of three hybrid systems:

  - Apple Mac operating system (laptop)

  - iOS (mobile operating systems)

  - Android (mobile operating systems)

Architecture of Mac OS X and iOS
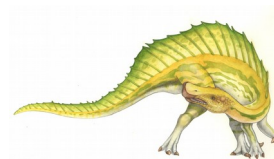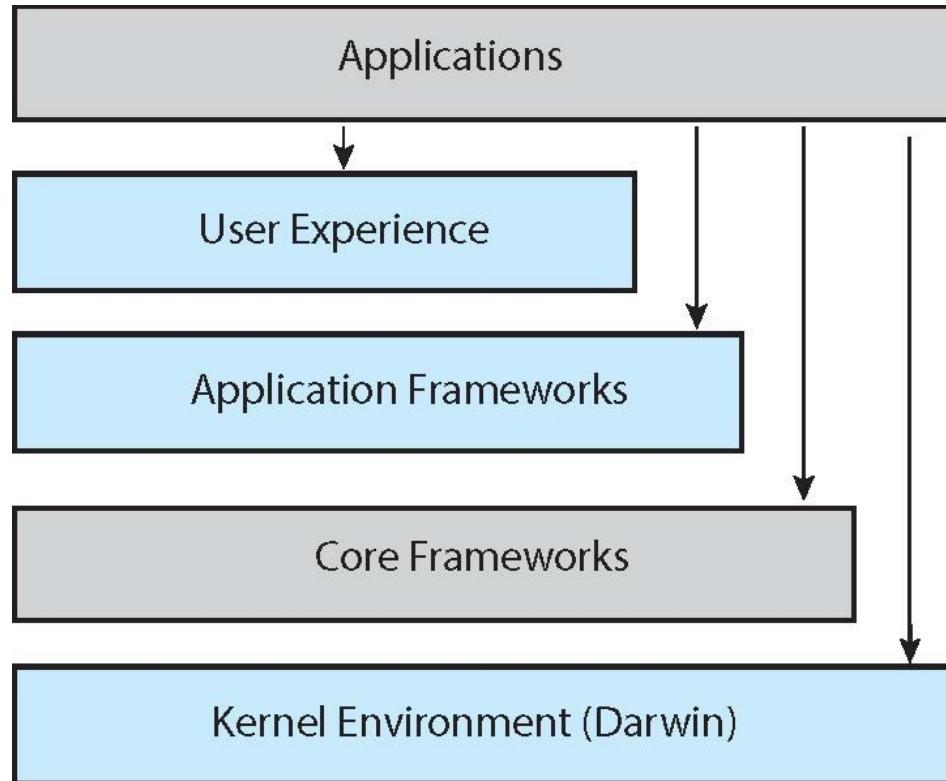
Darwin

- Darwin uses a hybrid structure, and is shown Darwin is a layered system which consists primarily of the Mach microkernel and the BSD UNIX kernel.
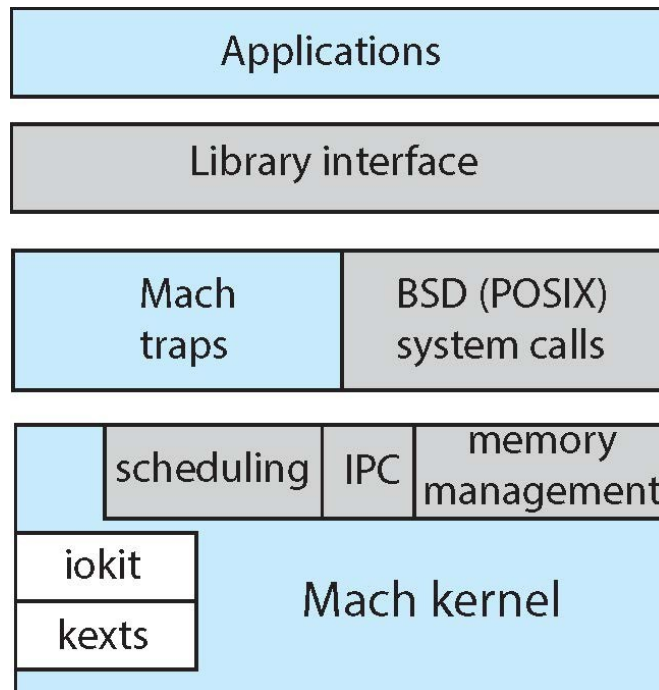
Android

- Developed by Open Handset Alliance (mostly Google)
  - Open Source
- Similar stack to iOS
- Based on Linux kernel but modified
  - Provides process, memory, device-driver management
  - Adds power management
- Runtime environment includes core set of libraries and Dalvik virtual machine
  - Apps developed in Java plus Android API
    - Java class files compiled to Java bytecode then translated to executable than runs in Dalvik VM
- Libraries include frameworks for web browser (webkit), database (SQLite), multimedia, smaller libc

Android Architecture

Applications

Application Framework

Libraries

SQLite	openGL
surface manager	media framework
webkit	libc

Android runtime

Core Libraries

Dalvik virtual machine

Linux kernel
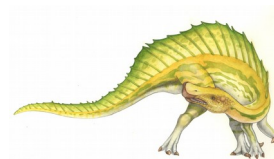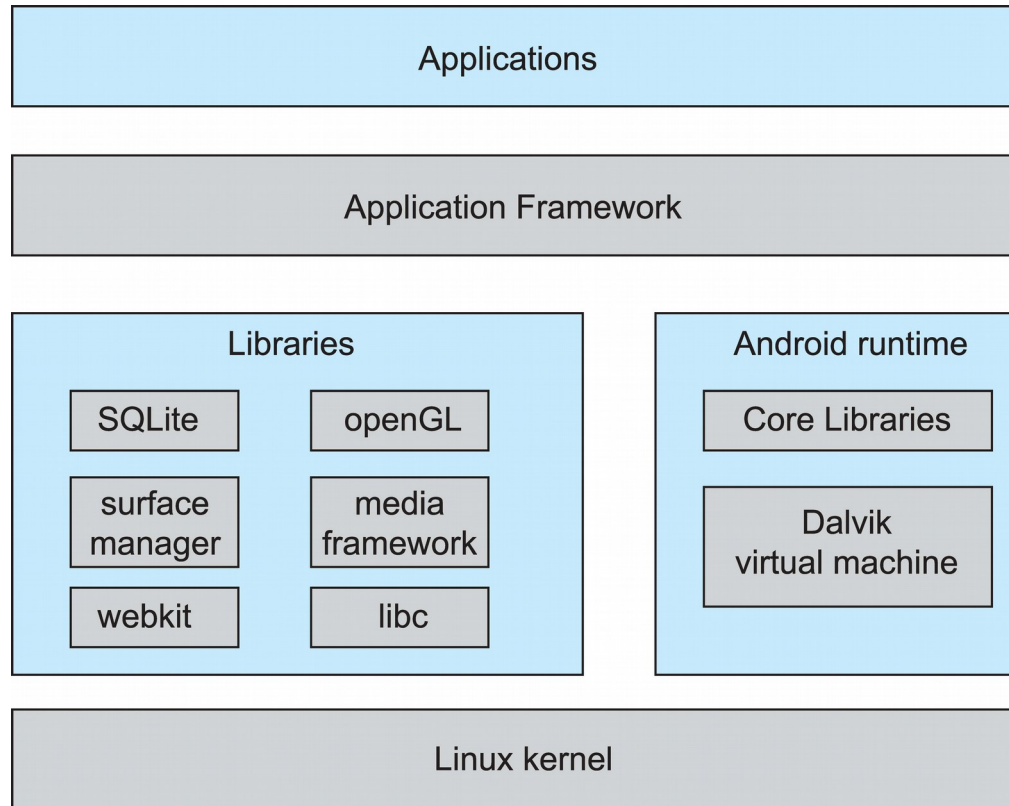
Operating-System Debugging

- **Debugging** is finding and fixing errors, or **bugs**
- OS generate **log files** containing error information
- Failure of an application can generate **core dump** file capturing memory of the process
- Operating system failure can generate **crash dump** file containing kernel memory
- Beyond crashes, performance tuning can optimize system performance
  - Sometimes using *trace listings* of activities, recorded for analysis
  - **Profiling** is periodic sampling of instruction pointer to look for statistical trends

Kernighan's Law: Debugging is twice as hard as writing the code in the first place. Therefore, if you write the code as cleverly as possible, you are, by definition, not smart enough to debug it."

Performance Tuning

- Improve performance by removing bottlenecks

- OS must provide means of computing and displaying measures of system behavior

- For example, "top" program or Windows Task Manager

DTrace

- DTrace tool in Solaris, FreeBSD, Mac OS X allows live instrumentation on production systems

- **Probes** fire when code is executed within a **provider**, capturing state data and sending it to **consumers** of those probes

- Example of following XEventsQueued system call move from libc library to kernel and back

```
# ./all.d 'pgrep xclock' XEventsQueued
dtrace: script './all.d' matched 52377 probes
CPU FUNCTION
  0  -> XEventsQueued                        U
  0    -> _XEventsQueued                     U
  0      -> _X11TransBytesReadable           U
  0      <- _X11TransBytesReadable           U
  0      -> _X11TransSocketBytesReadable     U
  0      <- _X11TransSocketBytesreadable     U
  0      -> ioctl                            U
  0        -> ioctl                          K
  0          -> getf                         K
  0            -> set_active_fd              K
  0            <- set_active_fd              K
  0          <- getf                         K
  0          -> get_udatamodel               K
  0          <- get_udatamodel               K
  ...
  0          -> releasef                     K
  0            -> clear_active_fd            K
  0            <- clear_active_fd            K
  0            -> cv_broadcast               K
  0            <- cv_broadcast               K
  0          <- releasef                     K
  0        <- ioctl                          K
  0      <- ioctl                            U
  0    <- _XEventsQueued                     U
  0  <- XEventsQueued                        U
```

- DTrace code to record amount of time each process with UserID 101 is in running mode (on CPU) in nanoseconds

```
sched:::on-cpu
uid == 101
{
   self->ts = timestamp;
}

sched:::off-cpu
self->ts
{
   @time[execname] = sum(timestamp - self->ts);
   self->ts = 0;
}
```

```
# dtrace -s sched.d
dtrace: script 'sched.d' matched 6 probes
^C
         gnome-settings-d          142354
         gnome-vfs-daemon          158243
         dsdm                      189804
         wnck-applet               200030
         gnome-panel               277864
         clock-applet              374916
         mapping-daemon            385475
         xscreensaver              514177
         metacity                  539281
         Xorg                      2579646
         gnome-terminal            5007269
         mixer_applet2             7388447
         java                      10769137
```

**Figure 2.21** Output of the D code.

Operating System Generation

- Operating systems are designed to run on any of a class of machines; the system must be configured for each specific computer site

- **SYSGEN** program obtains information concerning the specific configuration of the hardware system

  - Used to build system-specific compiled kernel or system-tuned

  - Can general more efficient code than one general kernel
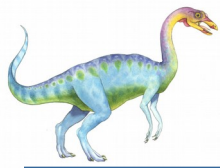
System Boot

- When power initialized on system, execution starts at a fixed memory location

  - Firmware ROM used to hold initial boot code

- Operating system must be made available to hardware so hardware can start it

  - Small piece of code – **bootstrap loader**, stored in **ROM** or **EEPROM** locates the kernel, loads it into memory, and starts it

  - Sometimes two-step process where **boot block** at fixed location loaded by ROM code, which loads bootstrap loader from disk

- Common bootstrap loader, **GRUB**, allows selection of kernel from multiple disks, versions, kernel options

- Kernel loads and system is then **running**

End of Chapter 2